

Lab #5 – Photon mapping (part 2)

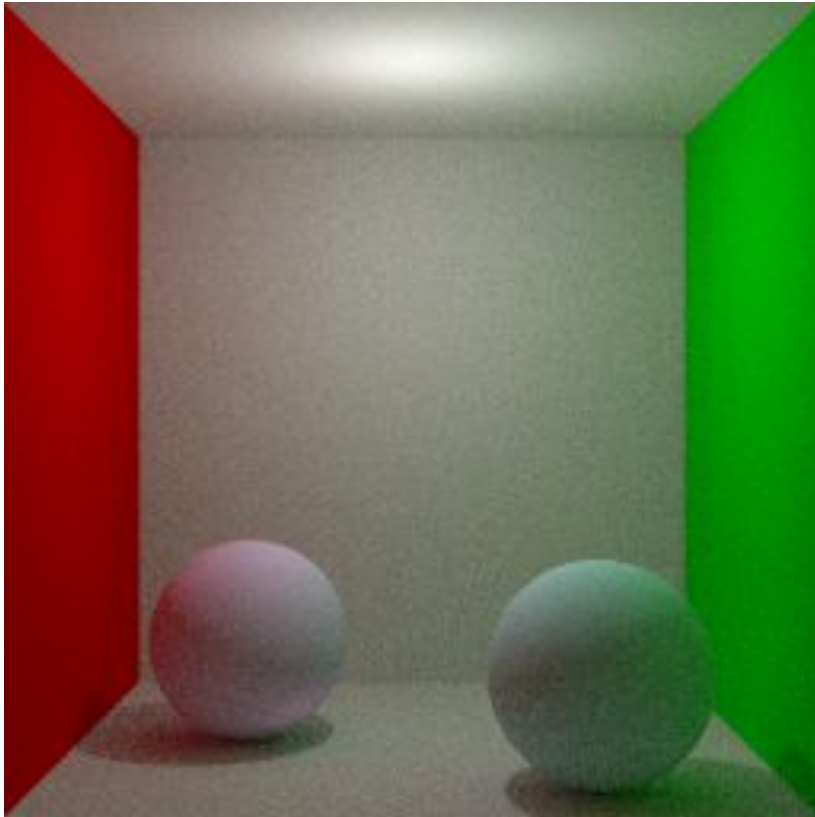
Informática Gráfica

Adolfo Muñoz - Julio Marco
Pablo Luesia - J. Daniel Subías – Óscar Pueyo




Before we begin...

- Today: render the first image with photon mapping



Path tracing




Photon mapping

Before we begin...


- Lab 5 (photon mapping) **is the second submitted work**
 - Recommended deadline: December 4th
 - Moodle: January 11th
- You can probably **reuse most of your code** for this assignment
- Remember: Final work is 80% of the final grade

Recap: photon mapping basics

- Photon mapping is a **two-pass** algorithm
 - Uses an **intermediate storage** known as the **photon map**




Lab 4 (path tracing)




Lab 5 (photon mapping)

Recap: photon map construction

- Split a path into two:
 - (1) lights write into the photon map 
 - (2) camera reads from the photon map
- Photons are sent out from the light sources (**photon random walks**)

Final photon map



Recap: photon map construction

The origin of the photons is the light source

- **Conservation of energy:**


$$p_1 + p_2 + p_3 + p_4 + p_5 + p_6 = 4\pi p_0$$

The light source emitted S photons.

A **photon's flux** is $p_i = 4\pi p_0 / S$


Careful with the photon map's size limit

You might need to update S after doing the random walks




Recap: how the photon map is used


- Split a path into two:
 - (1) lights write into the photon map
 - (2) camera reads from the photon map



$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$

Recap: how the photon map is used


- Split a path into two:
 - (1) lights write into the photon map
 - (2) camera reads from the photon map 




$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$

- Idea: use **k nearby** photons (distance < r) to approximate the Render Equation

Recap: how the photon map is used


- Split a path into two:
 - (1) lights write into the photon map
 - (2) camera reads from the photon map 




$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$

- Idea: use **k nearby** photons (distance < r) to approximate the Render Equation

Recap: how the photon map is used

- Split a path into two:
 - (1) lights write into the photon map
 - (2) camera reads from the photon map 




$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$

- Idea: use **k nearby** photons (distance $< r$) to approximate the Render Equation
- Integral is approximated as the **sum of k = 3 contributions**


Which color do we fill each pixel with?

- For path tracing:




$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$

The full integral




Which color do we fill each pixel with?

- For photon mapping:




Which color do we fill each pixel with?

- For photon mapping:
1) Calculate the photon map
(previous session)




Which color do we fill each pixel with?

- For photon mapping:
 - 1) Launch rays from the camera (**same as path tracing**)
 - 2) Intersect at point \mathbf{x}
 - 3) Search for nearest neighbours of \mathbf{x} in the photon map:




Which color do we fill each pixel with?

- For photon mapping:
 - 2) Launch rays from the camera (**same as path tracing**)
- Intersect at point x
- Search for nearest neighbours of x in the photon map:
 - a) k nearest photons
 - b) photons at distance $\leq r$



Which color do we fill each pixel with?

- For photon mapping:
 - 2) Launch rays from the camera (**same as path tracing**)
- Intersect at point x
- Search for nearest neighbours of x in the photon map:
 - a) k nearest photons
 - b) photons at distance $\leq r$
- Note how some photons are from **other surfaces**
(you can try to solve this, or not)





Which color do we fill each pixel with?

- For photon mapping:

3) Estimate color of the pixel with nearby photons

$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$




Kernel density estimation

$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$

Each photon p contains:

- Φ_p flux
- \mathbf{x}_p position
- ω_p direction



Kernel density estimation


$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$

Each photon p contains:

- Φ_p flux
- \mathbf{x}_p position
- ω_p direction

Constant density estimation: box kernel

$$L_o(\mathbf{x}, \omega_o) \approx \sum_{p=1}^k f_r(\mathbf{x}, \omega_p, \omega_o) \frac{\Phi_p}{\pi r_k^2}$$



Kernel density estimation

$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$


Each photon p contains:

- Φ_p flux
- \mathbf{x}_p position
- ω_p direction

Constant density estimation: box kernel

$$L_o(\mathbf{x}, \omega_o) \approx \sum_{p=1}^k f_r(\mathbf{x}, \omega_p, \omega_o) \frac{\Phi_p}{\pi r_k^2}$$

area of the green circle
 r_k is the radius that contains all photons



Kernel density estimation

$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$


Each photon p contains:

- Φ_p flux
- \mathbf{x}_p position
- ω_p direction

Far away from the
intersection point

Constant density estimation: box kernel

$$L_o(\mathbf{x}, \omega_o) \approx \sum_{p=1}^k f_r(\mathbf{x}, \omega_p, \omega_o) \frac{\Phi_p}{\pi r_k^2}$$



Kernel density estimation

$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$


Each photon p contains:

- Φ_p flux
- \mathbf{x}_p position
- ω_p direction

Non-constant density estimation (optional)

Gives more weight to photons closer to \mathbf{x}

$$L_o(\mathbf{x}, \omega_o) \approx \sum_{p=1}^k f_r(\mathbf{x}, \omega_p, \omega_o) \Phi_p \underline{K_{2D}(|\mathbf{x} - \mathbf{x}_p|, r_k)}$$



Kernel density estimation

$$L_o(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int_{\Omega} L_i(\mathbf{x}, \omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) |\mathbf{n} \cdot \omega_i| d\omega_i$$


Each photon p contains:

- Φ_p flux
- \mathbf{x}_p position
- ω_p direction

Non-constant density estimation (optional)

Gives more weight to photons closer to \mathbf{x}


$$L_o(\mathbf{x}, \omega_o) \approx \sum_{p=1}^k f_r(\mathbf{x}, \omega_p, \omega_o) \Phi_p K_{2D}(\|\mathbf{x} - \mathbf{x}_p\|, r_k)$$



Kernel density estimation


Constant density estimation

Box kernel




Non-constant density estimation


Cone kernel




Gaussian kernel




Photon mapping is biased




Photon mapping is biased




Example: area with hard shadows




Photon mapping is biased



Example: area with hard shadows




Photon mapping is biased



Example: area with hard shadows

When a ray intersects, it picks up photons from both areas

Photon mapping is biased




Example: area with hard shadows

When a ray intersects, it picks up photons from both areas

Produces a **biased estimation**


Photon mapping is biased



Example: area with hard shadows


Estimated result is **biased**

(This did not happen for path tracing)



How to compute direct light

We can compute direct light
without using a photon map:



How to compute direct light

We can compute direct light
without using a photon map:

However


You are duplicating energy



Next-event estimation



First bounce photons




How to compute direct light

We can compute direct light
without using a photon map:

However


You are duplicating energy



Next-event estimation

First bounce photons


Only choose one method, not both



How to compute direct light

Method 1:


- When generating the photon map, store photons on every diffuse interaction, including the first one



How to compute direct light

Method 1:


- When generating the photon map, store photons on every diffuse interaction, including the first one
- Later, compute direct light using kernel density estimation



How to compute direct light

Method 2:


- When generating the photon map, do not store the first bounce (direct) photons



How to compute direct light


Method 2:

- When generating the photon map, do not store the first bounce (direct) photons
- Later, compute direct light using next-event estimation and indirect light with kernel density estimation



Example scene: Cornell Box

- **Geometry**



Planes defined by normal (n) and distance (d)

Left plane $n = (1, 0, 0)$, $d = 1$

Right plane $n = (-1, 0, 0)$, $d = 1$

Floor plane $n = (0, 1, 0)$, $d = 1$

Ceiling plane $n = (0, -1, 0)$, $d = 1$

Back plane $n = (0, 0, -1)$, $d = 1$

Spheres defined by center (c) and radius (r)

Left sphere $c = (-0.5, -0.7, 0.25)$, $r = 0.3$


- Diffuse

Right sphere $c = (0.5, -0.7, -0.25)$, $r = 0.3$

- Diffuse

Example scene: Cornell Box

- Camera & light sources



Camera and image plane defined by

Origin $O = (0, 0, -3.5)$

Left $L = (-1, 0, 0)$


Up $U = (0, 1, 0)$

Forward $F = (0, 0, 3)$

Size 256x256 pixels

Example scene: Cornell Box

- Light sources



Center and power (emission)

Center $c = (0, 0.5, 0)$


Power can be any number e.g. $p = (1, 1, 1)$

Just be careful with the #MAX



```
1  P3
2  # feep.ppm
3  #MAX=<maximum of your RGB memory values>
4  4 4
5  15
6  0 0 0 0 0 0 0 0 0 0 15 0 15
```


Example scene: Cornell Box

- Results (direct light is computed using next-event estimation)



Using a point light



With tone mapping

DO ASK questions, either now or after the lab

But be reasonable, please :)

pluesia@unizar.es | dsubias@unizar.es | o.pueyo@unizar.es

What to expect from this session

In the programming language of your choice implement:

- Generate an image from your photon map:
 - Launch rays from the camera, intersect at point \mathbf{x} , find nearest photons
 - Use constant density estimation (box kernel) to estimate the Render Equation
 - Careful with direct light calculation!
- Recommended deadline: November 27th (moodle: January 11th)
 - Extensions (do not count towards recommended deadline):
 - **Recommended to finish base photon mapper before any optionals**
 - Try **different kernels** (cone, Gaussian, or more sophisticated ones)
 - Use an **adaptive kernel bandwidth** (radius)
 - Others: participating media, transient photon mapping, etc. (talk with us before)